

Structured Knowledge Accumulation: An Autonomous Framework for Layer-Wise Entropy Reduction in Neural Learning

Bouarfa Mahi
Quantiota
Email: info@quantiota.org

February 20, 2025
Version 1.8

Abstract

We introduce the Structured Knowledge Accumulation (SKA) framework, which redefines entropy as a dynamic, layer-wise measure of knowledge alignment in neural networks: $H^{(l)} = -\frac{1}{\ln 2} \sum_{k=1}^K \mathbf{z}_k^{(l)} \cdot \Delta \mathbf{D}_k^{(l)}$, where $\mathbf{z}_k^{(l)}$ is the knowledge vector and $\Delta \mathbf{D}_k^{(l)}$ is the decision probability shift vector at layer l over K forward steps. Rooted in the continuous entropy formulation $H = -\frac{1}{\ln 2} \int z dD$, SKA derives the sigmoid function, $D_i^{(l)} = \frac{1}{1+e^{-z_i^{(l)}}}$, as an emergent property of entropy minimization. This approach generalizes to fully connected networks without backpropagation, with each layer optimizing locally by aligning $\mathbf{z}_k^{(l)}$ with $\Delta \mathbf{D}_k^{(l)}$, guided by $\cos(\theta_k^{(l)})$. Total network entropy, $H = \sum_{l=1}^L H^{(l)}$, decreases hierarchically as knowledge structures evolve. Offering a scalable, biologically plausible alternative to gradient-based training, SKA bridges information theory and artificial intelligence, with potential applications in resource-constrained and parallel computing environments.

1 Introduction

Entropy, classically defined by Shannon [1] as $H = -\sum p_i \log_2 p_i$, quantifies uncertainty in a static, discrete probabilistic system. While foundational, this formulation falls short of capturing the dynamic, continuous structuring of knowledge in intelligent systems like neural networks. The sigmoid function, $\sigma(z) = \frac{1}{1+e^{-z}}$, a cornerstone of artificial intelligence (AI), has lacked a theoretical basis beyond empirical utility. Conventional training via backpropagation, which propagates errors backward through the network, is computationally intensive and biologically implausible, constraining scalability and real-world applicability.

This article presents the Structured Knowledge Accumulation (SKA) framework, reimagining entropy as a continuous process of knowledge accumulation. We propose:

1. Entropy as a dynamic measure, expressed in layers as $H^{(l)} = -\frac{1}{\ln 2} \sum_{k=1}^K \mathbf{z}_k^{(l)} \cdot \Delta \mathbf{D}_k^{(l)}$, approximating the continuous $H = -\frac{1}{\ln 2} \int z dD$,
2. Structured knowledge (z) and its continuous accumulation as the foundation of learning,
3. A forward-only, backpropagation-free learning rule driven by local entropy minimization.

We derive the sigmoid from continuous entropy reduction, extend it to fully connected networks, and demonstrate learning through local alignment of knowledge with decision dynamics. This framework offers a novel perspective for AI, enhancing optimization efficiency, interpretability, and biological plausibility, with implications for scalable and distributed neural systems.

2 Redefining Entropy in the SKA Framework

Entropy traditionally quantifies uncertainty in probabilistic systems, but its classical form is static and discrete, limiting its applicability to dynamic learning processes like those in neural networks. In the Structured Knowledge Accumulation (SKA) framework, we redefine entropy as a continuous, evolving measure that reflects knowledge alignment over time or processing steps. This section contrasts Shannon’s discrete entropy with our continuous reformulation, enabling the use of continuous decision probabilities and supporting the derivation of the sigmoid function through entropy minimization.

2.1 Classical Shannon Entropy

For a binary system with decision probability D , Shannon’s entropy is:

$$H = -D \log_2 D - (1 - D) \log_2 (1 - D) \quad (1)$$

Its derivative with respect to D is:

$$\frac{dH}{dD} = \log_2 \left(\frac{1 - D}{D} \right) \quad (2)$$

This formulation assumes D is a fixed probability, typically associated with discrete outcomes (e.g., 0 or 1). While foundational, it does not capture the continuous evolution of knowledge in a learning system, where D may vary smoothly as the network processes inputs. To address this, we seek a continuous entropy measure that accommodates dynamic changes in D , aligning with the SKA’s focus on knowledge accumulation.

2.2 Entropy as Knowledge Accumulation

In SKA, we redefine entropy for a single neuron as a continuous process:

$$H = -\frac{1}{\ln 2} \int z dD \quad (3)$$

Here, z represents the neuron’s structured knowledge, and dD is an infinitesimal change in the decision probability, treated as a continuous variable over the range $[0, 1]$. The factor $-\frac{1}{\ln 2}$ ensures alignment with base-2 logarithms, consistent with Shannon’s information units. Unlike the static snapshot of Equation 1, this integral captures how entropy accumulates as z drives changes in D , reflecting a dynamic learning process.

For a layer l with n_l neurons over K forward steps, we approximate this continuous form discretely:

$$H^{(l)} = -\frac{1}{\ln 2} \sum_{k=1}^K \mathbf{z}_k^{(l)} \cdot \Delta \mathbf{D}_k^{(l)} \quad (4)$$

where $\mathbf{z}_k^{(l)} = [z_1^{(l)}(k), \dots, z_{n_l}^{(l)}(k)]^T$ is the knowledge vector, $\Delta \mathbf{D}_k^{(l)} = [\Delta D_1^{(l)}(k), \dots, \Delta D_{n_l}^{(l)}(k)]^T$ is the vector of decision probability shifts, and the scalar product is:

$$\mathbf{z}_k^{(l)} \cdot \Delta \mathbf{D}_k^{(l)} = \sum_{i=1}^{n_l} z_i^{(l)}(k) \Delta D_i^{(l)}(k) \quad (5)$$

The total network entropy sums over all layers:

$$H = \sum_{l=1}^L H^{(l)} \quad (6)$$

Equation 3 is the core theoretical construct, with Equation 4 as its practical discrete approximation. As $K \rightarrow \infty$ and $\Delta \mathbf{D}_k^{(l)}$ becomes infinitesimally small, Equation 4 approaches the continuous integral, enabling us to model D as a smooth function of z . This continuous perspective is essential for deriving the sigmoid using dynamics in later sections, while the discrete form supports implementation in neural architectures.

2.3 Accumulated Knowledge

Knowledge accumulates over steps:

$$z_k = z_0 + \sum_{f=1}^k \Delta z_f \quad (7)$$

In a layer, $\mathbf{z}_k^{(l)}$ evolves, reducing $H^{(l)}$ as it aligns with $\Delta \mathbf{D}_k^{(l)}$.

3 Deriving the Sigmoid Function

The SKA framework posits that the sigmoid function emerges naturally from continuous entropy minimization, linking structured knowledge to decision probabilities. This section demonstrates that when D follows $D = \frac{1}{1+e^{-z}}$, the SKA entropy H_{SKA} equals the classical Shannon entropy H_{Shannon} , differing by a constant (zero). By leveraging the continuous formulation from Section 2, we derive this equivalence, reinforcing the framework's theoretical grounding.

3.1 Key Definitions

3.1.1 Shannon Entropy (for binary decisions)

For a binary system with continuous decision probability D :

$$H_{\text{Shannon}} = -D \log_2 D - (1 - D) \log_2 (1 - D) \quad (8)$$

3.1.2 SKA Entropy (layer-wise, for a single neuron)

The SKA entropy, defined continuously as in Section 2.2, is:

$$H_{\text{SKA}} = -\frac{1}{\ln 2} \int z dD, \quad (9)$$

where $z = -\ln\left(\frac{1-D}{D}\right)$ relates knowledge to D , consistent with $D = \frac{1}{1+e^{-z}}$ as shown in Section 3.1.

3.2 Equivalence Proof

Substituting $z = -\ln\left(\frac{1-D}{D}\right)$ (or equivalently, $z = \ln\left(\frac{D}{1-D}\right)$) into H_{SKA} :

$$H_{\text{SKA}} = -\frac{1}{\ln 2} \int \ln\left(\frac{D}{1-D}\right) dD. \quad (10)$$

Evaluate the integral with substitution $u = D$, $du = dD$:

$$\int \ln\left(\frac{D}{1-D}\right) dD = D \ln\left(\frac{D}{1-D}\right) + \ln(1-D). \quad (11)$$

Substituting back:

$$H_{\text{SKA}} = -\frac{1}{\ln 2} \left[D \ln \left(\frac{D}{1-D} \right) + \ln(1-D) \right]. \quad (12)$$

Rewrite $\ln \left(\frac{D}{1-D} \right) = \ln D - \ln(1-D)$:

$$H_{\text{SKA}} = -\frac{1}{\ln 2} [D \ln D - D \ln(1-D) + \ln(1-D)]. \quad (13)$$

Factorize:

$$H_{\text{SKA}} = -\frac{1}{\ln 2} [D \ln D + (1-D) \ln(1-D)]. \quad (14)$$

Thus:

$$H_{\text{SKA}} = H_{\text{Shannon}}. \quad (15)$$

3.3 Implications

- **Zero Difference:** The SKA and Shannon entropies are identical (differing by zero) when $D = \frac{1}{1+e^{-z}}$, confirming the sigmoid as an emergent property of continuous entropy reduction.
- **Knowledge Alignment:** This equivalence stems from z structuring D to minimize uncertainty, as defined in Sections 2 and 3.

3.4 Significance

1. **Theoretical Consistency:** SKA extends Shannon entropy into a continuous, dynamic context while preserving its core properties for sigmoidal outputs.
2. **Backpropagation-Free Learning:** Since $H_{\text{SKA}} = H_{\text{Shannon}}$, layer-wise entropy minimization aligns with classical uncertainty reduction, achieved via forward dynamics alone.
3. **Biological Plausibility:** The continuous, local alignment of z with D mirrors plausible neural learning mechanisms.

3.5 Summary

When D is the sigmoid function, H_{SKA} matches H_{Shannon} exactly, with a difference of zero. This result, derived from the continuous entropy $H_{\text{SKA}} = -\frac{1}{\ln 2} \int z dD$, validates SKA's foundation and its seamless integration with information theory, leveraging continuous dynamics for neural learning with classical information theory.

4 The Fundamental Law of Entropy Reduction

The SKA framework establishes a fundamental law governing how entropy decreases as structured knowledge evolves. This section derives this law using continuous dynamics, reflecting the continuous nature of decision probabilities and knowledge accumulation introduced in Sections 2 and 3. We then provide a discrete approximation for practical implementation, ensuring the framework's applicability to neural networks while rooting it in a continuous theoretical foundation.

4.1 Continuous Dynamics

For a single neuron, the rate of entropy change with respect to structured knowledge z is derived from the continuous entropy $H = -\frac{1}{\ln 2} \int z dD$. Taking the partial derivative:

$$\frac{\partial H}{\partial z} = -\frac{1}{\ln 2} z D (1 - D) \quad (16)$$

This follows from $D = \frac{1}{1+e^{-z}}$ (as derived in Section 4), where $\frac{dD}{dz} = D(1 - D)$, and reflects the neuron's local contribution to entropy reduction. For a layer l with n_l neurons at step k , this extends to each neuron i :

$$\frac{\partial H^{(l)}}{\partial z_i^{(l)}(k)} = -\frac{1}{\ln 2} z_i^{(l)}(k) D_i^{(l)}(k) (1 - D_i^{(l)}(k)) \quad (17)$$

Equation 17 governs the continuous reduction of layer-wise entropy $H^{(l)}$, driven by the alignment of $z_i^{(l)}(k)$ with the sigmoidal decision probability $D_i^{(l)}(k)$. This dynamic, localized process underpins the SKA's forward-only learning mechanism, leveraging the continuous evolution of D over time or input processing.

4.2 Discrete Dynamics

In practice, neural networks operate over discrete forward steps. For a single neuron at step k , the entropy gradient approximates the continuous form, incorporating the change in decision probability $\Delta D_k = D_k - D_{k-1}$:

$$\left. \frac{\partial H}{\partial z} \right|_k = -\frac{1}{\ln 2} [z_k D_k (1 - D_k) + \Delta D_k] \quad (18)$$

For layer l at step k , this becomes:

$$\frac{\partial H^{(l)}}{\partial z_i^{(l)}(k)} = -\frac{1}{\ln 2} z_i^{(l)}(k) \left[D_i^{(l)}(k) (1 - D_i^{(l)}(k)) + \Delta D_i^{(l)}(k) \right] \quad (19)$$

Equation 19 adapts the continuous law to discrete steps, where $\Delta D_i^{(l)}(k)$ captures the step-wise shift in $D_i^{(l)}(k)$. While Equation 17 represents the ideal continuous dynamics, Equation 19 provides a computable approximation, aligning knowledge adjustments with observed changes in decision probabilities across discrete iterations.

5 Generalization to Fully Connected Networks

The SKA framework extends seamlessly from single neurons to fully connected neural networks, leveraging the continuous entropy reduction principles established earlier. For a network with L layers, knowledge and decision probabilities evolve hierarchically, reducing total entropy through local, forward-only adjustments. This section outlines how SKA operates across layers, maintaining its biologically plausible and scalable design.

For a network with L layers:

- $\mathbf{z}_k^{(l)} = \mathbf{W}^{(l)} \mathbf{x}_k^{(l-1)} + \mathbf{b}^{(l)}$, the knowledge vector at layer l and step k ,
- $\mathbf{D}_k^{(l)} = \sigma(\mathbf{z}_k^{(l)})$, the decision probabilities derived via the sigmoid function,
- $\Delta \mathbf{D}_k^{(l)} = \mathbf{D}_k^{(l)} - \mathbf{D}_{k-1}^{(l)}$, the step-wise shift in decision probabilities.

Layer-wise entropy, rooted in the continuous formulation, is approximated discretely:

$$H^{(l)} = -\frac{1}{\ln 2} \sum_{k=1}^K \mathbf{z}_k^{(l)} \cdot \Delta \mathbf{D}_k^{(l)} \quad (20)$$

The alignment between knowledge and decision shifts is quantified as:

$$\mathbf{z}_k^{(l)} \cdot \Delta \mathbf{D}_k^{(l)} = \|\mathbf{z}_k^{(l)}\| \|\Delta \mathbf{D}_k^{(l)}\| \cos(\theta_k^{(l)}) \quad (21)$$

Total network entropy aggregates across layers:

$$H = \sum_{l=1}^L H^{(l)} \quad (22)$$

Learning proceeds by aligning $\mathbf{z}_k^{(l)}$ with $\Delta \mathbf{D}_k^{(l)}$ at each layer, reducing $H^{(l)}$ locally without requiring backward error propagation. In the continuous limit, this alignment reflects a smooth evolution of knowledge, approximated here by discrete steps for computational feasibility.

6 Learning Without Backpropagation

SKA achieves learning through localized entropy minimization, eliminating the need for backpropagation by leveraging forward-only dynamics. This section details the weight update mechanism and supporting metrics, grounded in the continuous entropy reduction law, and adapted for discrete implementation in fully connected networks.

Entropy minimization at layer l is driven by:

$$\frac{\partial H^{(l)}}{\partial w_{ij}^{(l)}} = -\frac{1}{\ln 2} \sum_{k=1}^K \frac{\partial(\mathbf{z}_k^{(l)} \cdot \Delta \mathbf{D}_k^{(l)})}{\partial w_{ij}^{(l)}} \quad (23)$$

The update rule adjusts weights forward:

$$w_{ij}^{(l)} \leftarrow w_{ij}^{(l)} - \eta \frac{\partial H^{(l)}}{\partial w_{ij}^{(l)}} \quad (24)$$

Here, $\Delta D_i^{(l)}(k)$ is computed directly from forward passes, bypassing the need for error backpropagation. This local adjustment aligns with the continuous dynamics of knowledge evolution, approximated over discrete steps.

Step-wise Entropy Change

To quantify knowledge accumulation, the step-wise entropy change at layer l and step k is:

$$\Delta H_k^{(l)} = H_k^{(l)} - H_{k-1}^{(l)} = -\frac{1}{\ln 2} \mathbf{z}_k^{(l)} \cdot \Delta \mathbf{D}_k^{(l)} \quad (25)$$

This measures uncertainty reduction as $\mathbf{z}_k^{(l)}$ aligns with $\Delta \mathbf{D}_k^{(l)}$, with total layer entropy as:

$$H^{(l)} = \sum_{k=1}^K \Delta H_k^{(l)} \quad (26)$$

Entropy Gradient

The gradient of $H^{(l)}$ with respect to $\mathbf{z}_k^{(l)}$ at step k is:

$$\nabla H^{(l)} = \frac{\partial H^{(l)}}{\partial \mathbf{z}_k^{(l)}} = -\frac{1}{\ln 2} \mathbf{z}_k^{(l)} \odot \mathbf{D}_k'^{(l)} - \Delta \mathbf{D}_k^{(l)} \quad (27)$$

where $\mathbf{D}_k'^{(l)} = \mathbf{D}_k^{(l)} \odot (\mathbf{1} - \mathbf{D}_k^{(l)})$ is the sigmoid derivative. This gradient guides updates to minimize $H^{(l)}$, aligning knowledge with decision shifts.

Knowledge Evolution Across Layers

The gradient $\nabla H^{(l)} = -\frac{1}{\ln 2} \mathbf{z}_k^{(l)} \odot \mathbf{D}_k'^{(l)} + \Delta \mathbf{D}_k^{(l)}$ drives entropy reduction in layer l at step k . As $\mathbf{D}_k^{(l-1)}$ feeds into $\mathbf{z}_k^{(l)}$, each layer adapts uniquely—extracting broad features early and refining decisions later—mirroring a continuous knowledge flow approximated discretely.

Governing Equation of SKA

The network evolves according to:

$$\nabla H^{(l)} + \frac{1}{\ln 2} \mathbf{z}_k^{(l)} \odot \mathbf{D}_k'^{(l)} + \Delta \mathbf{D}_k^{(l)} = 0 \quad (28)$$

where $\nabla H^{(l)}$ minimizes entropy layer-wise, with updates following $-\nabla H^{(l)}$ to align $\mathbf{z}_k^{(l)}$ with $\Delta \mathbf{D}_k^{(l)}$.

Inter-Layer Entropy Change

The entropy change between layers l and $l+1$ at step k is:

$$\Delta H_k^{(l,l+1)} = -\frac{1}{\ln 2} \left[\mathbf{z}_k^{(l+1)} \cdot \Delta \mathbf{D}_k^{(l+1)} - \mathbf{z}_k^{(l)} \cdot \Delta \mathbf{D}_k^{(l)} \right] \quad (29)$$

This quantifies the spatial evolution of knowledge, complementing the temporal guidance of $\nabla H^{(l)}$, as entropy decreases through the network.

7 Application to Neural Networks

SKA structures knowledge hierarchically across layers, reducing total entropy H through continuous dynamics approximated over discrete steps. A multilayer perceptron (MLP) can train by minimizing H , with $\cos(\theta_k^{(l)})$ serving as a practical metric to monitor alignment between $\mathbf{z}_k^{(l)}$ and $\Delta \mathbf{D}_k^{(l)}$. This forward-only process leverages the framework’s scalability and autonomy, applicable to diverse network architectures.

Layer-wise Entropy Reduction in SKA

$$H = \sum_{l=1}^L H^{(l)} \downarrow \text{ as knowledge structures}$$

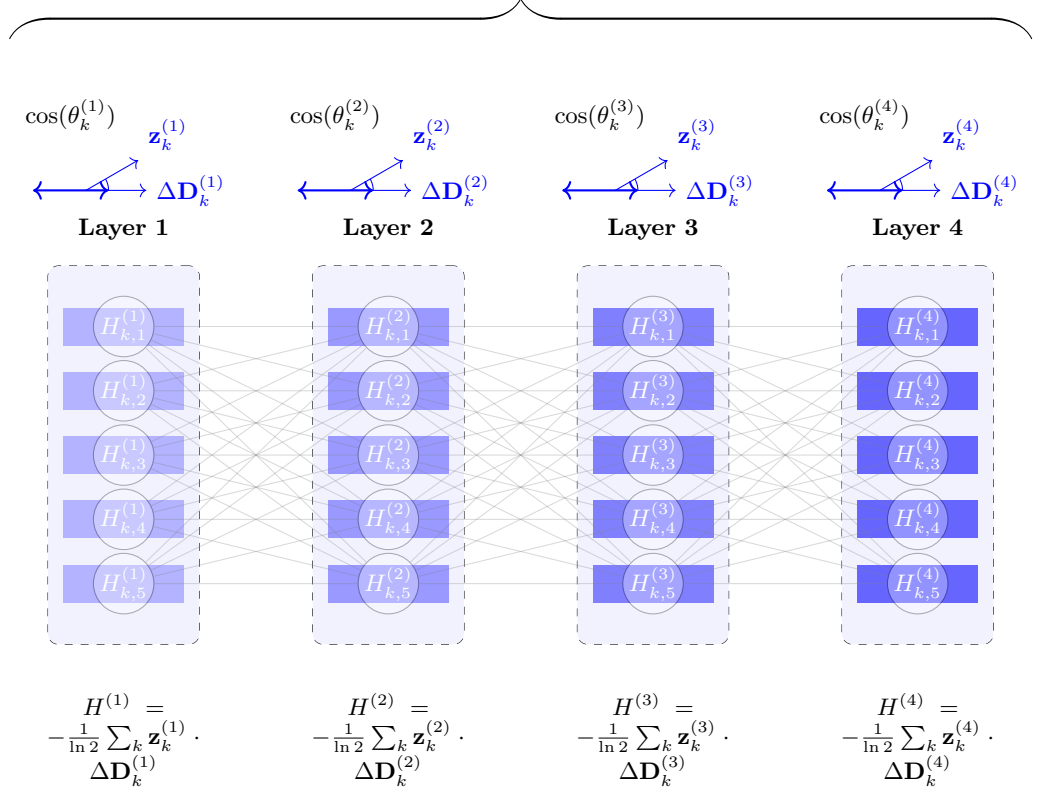


Figure 1: Layer-wise entropy reduction across layers at step k . Color intensity represents entropy level (darker blue = lower entropy), decreasing from Layer 1 to Layer 4. Each layer locally minimizes entropy by aligning knowledge vectors $\mathbf{z}_k^{(l)}$ with decision change vectors $\Delta \mathbf{D}_k^{(l)}$, measured by $\cos(\theta_k^{(l)})$.

8 Tensor-Based Implementation of SKA

To enhance computational efficiency and scalability, we introduce a **tensor-based implementation** of the Structured Knowledge Accumulation (SKA) framework. This approach preserves the theoretical foundations of SKA while leveraging tensor operations to enable efficient, parallelizable learning. By structuring SKA’s mathematical formulation within a tensor framework, we streamline the **computation of knowledge accumulation, entropy reduction, and decision updates** across multiple layers and time steps.

Tensor Definitions

A neural network operating under SKA can be represented using the following tensors:

- **Knowledge Tensor (\mathbf{Z}):** Encodes structured knowledge for each neuron across **layers** (L), **steps** (K), and **neurons** (n_{\max}).
- **Decision Probability Tensor (\mathbf{D}):** Stores the neuron activations as sigmoid-transformed knowledge values.

- **Shift Tensor ($\Delta\mathbf{D}$):** Captures **local probability shifts** between consecutive steps, reflecting SKA’s entropy-driven learning process.
- **Weight Tensor (\mathbf{W}) and Bias Tensor (\mathbf{b}):** Define learnable parameters that evolve dynamically to optimize knowledge accumulation.

Forward-Only Learning and Knowledge Update

SKA updates knowledge through a forward-only mechanism:

$$\mathbf{Z} = \mathbf{W} \cdot \mathbf{X} + \mathbf{b}$$

where \mathbf{X} is the input tensor. Unlike backpropagation, SKA **does not require gradients to propagate backward**, making it more computationally efficient.

Entropy Computation and Learning

SKA’s entropy formulation can be naturally expressed using tensor operations:

$$\mathbf{H} = -\frac{1}{\ln 2} \sum_{k=1}^{K-1} \mathbf{Z}_{k+1} \cdot \Delta\mathbf{D}_k$$

where **entropy decreases layer by layer** as structured knowledge accumulates. To optimize learning, entropy gradients are computed as:

$$\nabla\mathbf{H} = -\frac{1}{\ln 2} \mathbf{Z} \odot \mathbf{D}' - \Delta\mathbf{D}$$

where \mathbf{D}' is the derivative of the sigmoid function.

Weight Updates and Learning Stability

Weight updates follow a direct entropy minimization rule:

$$\mathbf{W} \leftarrow \mathbf{W} - \eta \frac{\partial\mathbf{H}}{\partial\mathbf{W}}$$

Additionally, an **alignment tensor** monitors whether knowledge updates follow expected trajectories:

$$\Theta_{l,k} = \cos(\theta_{k+1}^{(l)}) = \frac{\mathbf{Z}_{l,k+1} \cdot \Delta\mathbf{D}_{l,k}}{\|\mathbf{Z}_{l,k+1}\| \|\Delta\mathbf{D}_{l,k}\|}$$

This ensures SKA maintains a **structured** and **controlled** learning process.

Advantages of the Tensor-Based Implementation

- **Parallel Computation:** Tensor operations allow efficient computation across multiple neurons and layers.
- **Scalability:** The framework extends naturally to large-scale networks without increasing complexity.
- **Biologically Inspired Learning:** Forward-only updates make SKA suitable for real-time, **low-energy computing environments**.

This tensor-based formulation further strengthens SKA’s position as a **computationally efficient, biologically plausible, and scalable learning paradigm**. Future research will benchmark its performance against traditional backpropagation models, exploring its potential in edge computing, real-time AI, and neural network optimization.

9 Entropy Evolution in SKA

A key insight emerging from the SKA framework is the structured evolution of entropy across layers during learning. Unlike traditional deep networks, where entropy dynamics vary non-uniformly across layers, SKA exhibits a remarkable phenomenon: **all layers tend to converge toward a common entropy equilibrium value**. This behavior suggests that entropy is not merely decreasing but self-organizing in a structured manner.

9.1 Empirical Observation: Layer-Wise Entropy Convergence

Figure 2 presents the entropy evolution of different layers over multiple forward learning steps. The key observations are:

- **Layer-Specific Minima:** Each layer reaches a local minimum entropy value at different forward steps (K), before increasing slightly and stabilizing.
- **Convergence Toward a Common Equilibrium:** The entropy of layers 2, 3, and 4 asymptotically approaches a shared value around step $K = 49$, suggesting a fundamental equilibrium phenomenon.
- **Slow Convergence of Layer 1:** While layer 1 exhibits similar entropy reduction behavior, it approaches the equilibrium point more gradually, indicating a foundational role in hierarchical knowledge structuring.

This structured entropy evolution indicates that SKA does not merely reduce entropy but dynamically redistributes it across layers in a systematic fashion.




Figure 2: Entropy evolution across layers in an SKA neural network. Layers 2, 3, and 4 converge to a shared entropy equilibrium at $K = 49$, while Layer 1 approaches this point gradually.

9.2 Cosine Alignment Evolution

In addition to entropy evolution, SKA demonstrates another remarkable phenomenon: **cosine alignment evolution**. Figure 3 shows how the alignment between knowledge vectors $\mathbf{z}_k^{(l)}$ and decision probability shifts $\Delta \mathbf{D}_k^{(l)}$ evolves over time.

- **Cosine Alignment Convergence:** Across layers, $\cos(\theta_k^{(l)})$ converges toward a stable value, indicating increasing coherence between knowledge accumulation and decision updates.
- **Parallel Behavior to Entropy:** The stabilization of cosine values aligns with entropy reaching an equilibrium, reinforcing the structured learning dynamics of SKA.
- **Layer-Wise Synchronization:** The synchronization of alignment across layers suggests a self-organizing mechanism, optimizing knowledge updates through entropy minimization.

This empirical evidence further supports the hypothesis that SKA exhibits a natural, structured progression of learning, distinct from conventional gradient-based methods.




Figure 3: Cosine alignment evolution in SKA. The values of $\cos(\theta_k^{(l)})$ stabilize over forward steps, indicating structured knowledge alignment across layers.

The structured nature of entropy and cosine alignment evolution offers a novel perspective on learning in neural networks and may provide deeper insights into how biological and artificial learning systems self-organize their internal representations.

9.3 Output Decision Probability Evolution

Beyond entropy and cosine alignment, another critical visualization in SKA is the **evolution of output decision probabilities** over forward steps. Figure 4 illustrates the distribution of mean decision probabilities across all 10 classes as learning progresses.

- **Gradual Separation of Classes:** The decision probabilities shift over time, showing how SKA refines its class separability without explicit gradient updates.
- **Emergent Stability:** As the steps increase, decision probabilities stabilize, indicating convergence toward structured classification.
- **No Drastic Separation:** Unlike classical models, SKA does not create sharp decision boundaries but gradually refines knowledge accumulation to maintain a probabilistic structure.

This visualization provides further evidence of SKA’s **self-organizing, entropy-driven classification process**, reinforcing its fundamental distinction from traditional backpropagation-based learning.

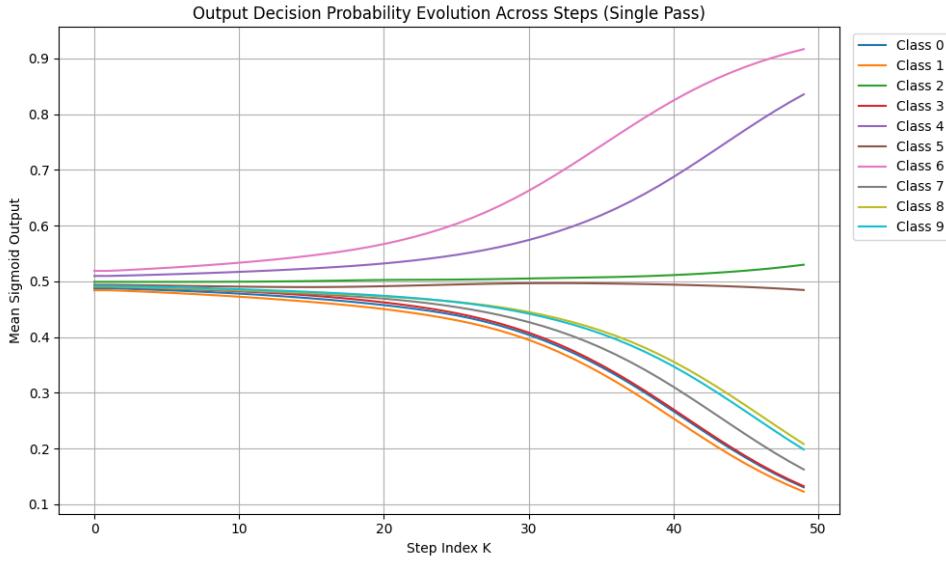


Figure 4: Evolution of output decision probabilities in SKA over forward steps. Unlike conventional models, SKA refines class separability progressively while maintaining structured probability distributions.

9.4 Frobenius Norm Evolution of the Knowledge Tensor

In addition to monitoring entropy, cosine alignment, and output decision probabilities, another illuminating perspective in SKA is the **Frobenius norm** of the knowledge tensor $\mathbf{z}^{(l)}$. For a layer l , the Frobenius norm is defined as

$$\|\mathbf{z}^{(l)}\|_F = \sqrt{\sum_{i,j} (z_{ij}^{(l)})^2},$$

where $z_{ij}^{(l)}$ is the knowledge value of neuron j in sample i (before the sigmoid activation).

- **Layer-Specific Magnitude Growth:** Each layer’s Frobenius norm reflects the overall magnitude of its knowledge tensor $\mathbf{z}^{(l)}$. A larger norm indicates that the pre-sigmoid activations are more extreme, suggesting stronger or more polarized responses. Interestingly, while some layers may exhibit rapid increases in their norms, our observations show that the final layer (Layer 4) tends to grow more slowly. This gradual increase in Layer 4’s Frobenius norm suggests that, despite its role in driving the output logits, its activations remain relatively moderate—possibly indicating an early stabilization of the output during the SKA learning process.

- **Single-Pass Dynamics:** Under a single-pass, forward-only scheme, some layers may exhibit steadily increasing norms, reflecting the absence of a backward error signal that would typically constrain large activations. This effect can be particularly pronounced in the final layer, where classification logits may grow larger as the model strives to minimize local entropy.
- **Relationship to Entropy and Alignment:** While entropy and cosine alignment measure how well the knowledge tensor $\mathbf{z}^{(l)}$ aligns with the decision shifts $\Delta\mathbf{D}^{(l)}$, the Frobenius norm focuses solely on the *magnitude* of $\mathbf{z}^{(l)}$. Thus, a layer may have a large norm yet still maintain low entropy if its knowledge vectors are well-aligned with the decision shifts.

Figure 5 illustrates how the Frobenius norms of the knowledge tensors evolve over multiple forward steps. Notably, layers can exhibit varying rates of growth or convergence, offering insights into how strongly each layer is “pushing” its logits to reduce local uncertainty.

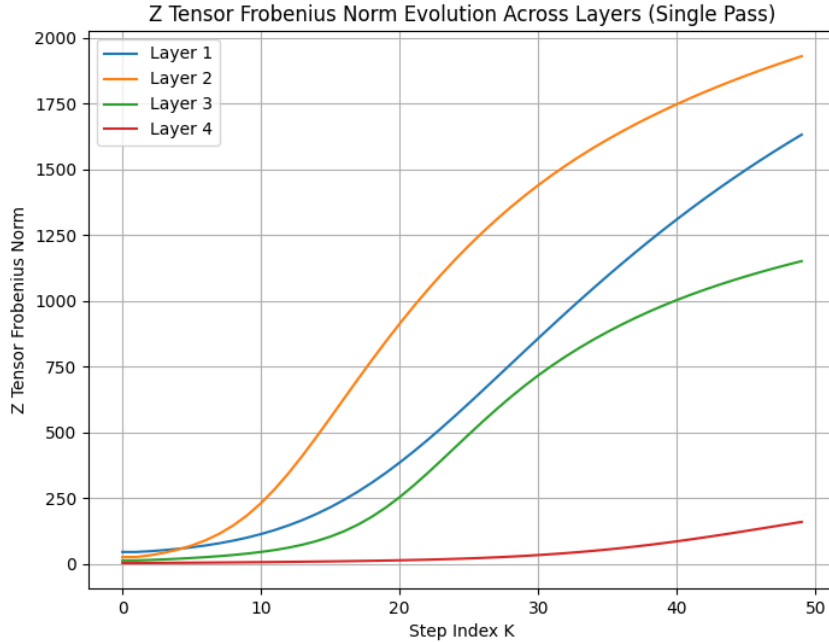


Figure 5: Frobenius norm evolution of the knowledge tensors $\mathbf{z}^{(l)}$ across layers during single-pass SKA training. Layers with larger norms typically exhibit stronger activations, reflecting the forward-only updates aimed at local entropy minimization.

Overall, tracking the Frobenius norm provides an additional lens through which to interpret **how much magnitude** each layer’s knowledge representation carries. Taken together with entropy, cosine alignment, and decision probability evolution, it completes a multi-faceted view of how SKA networks self-organize their internal representations in a purely forward-driven manner.

9.5 Entropy Trajectories

A crucial observation in SKA is the presence of structured entropy trajectories, particularly when plotted against the knowledge magnitude. Figure 6 illustrates the relationship between entropy reduction and the Frobenius norm of the knowledge tensor across layers.

- **U-Shaped Relationship:** Each layer exhibits a characteristic U-shaped curve, where entropy reduction initially decreases with increasing knowledge but later reverses after reaching a minimum point.
- **Progressive Shift of the Minimum:** The entropy minimum occurs at a progressively lower knowledge magnitude as we move from Layer 1 to Layer 4. This suggests that deeper layers require less overall knowledge magnitude to achieve entropy minimization.
- **Monotonic Knowledge Growth:** Unlike traditional models requiring explicit constraints or regularization, SKA naturally regulates knowledge accumulation without requiring manual tuning.

This structured behavior implies that entropy minimization is directly governed by the dynamics of knowledge accumulation, reinforcing the idea that SKA is inherently self-organizing.

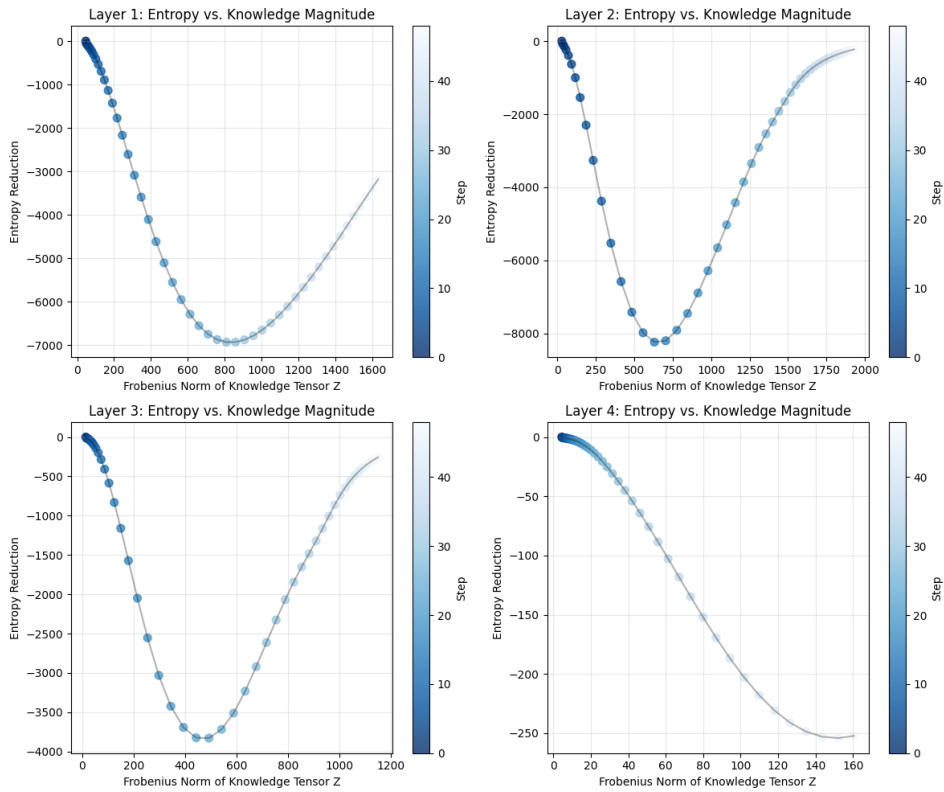


Figure 6: Entropy trajectories across layers, plotted against the Frobenius norm of the knowledge tensor. Each layer exhibits a U-shaped entropy curve, with the minimum shifting towards lower knowledge magnitudes as depth increases.

This observation provides a novel perspective on how entropy-driven learning occurs in SKA, demonstrating that each layer optimally structures knowledge without requiring external intervention.

9.6 The Tensor Net Function

A critical insight in the SKA framework is the **Tensor Net function**, which measures the balance between entropy and decision probability with respect to knowledge. For each layer l at step k , the Tensor Net function is defined as:

$$\text{Net}_k^{(l)} = \sum (\mathbf{D}_k^{(l)} - \mathbf{H}_k^{(l)}) \cdot \Delta \mathbf{Z}_k^{(l)} \quad (30)$$

where $\mathbf{D}_k^{(l)}$ represents the decision probability tensor, $\mathbf{H}_k^{(l)}$ is the entropy tensor, and $\Delta \mathbf{Z}_k^{(l)}$ is the change in knowledge tensor between steps. This function captures the relationship between decision-making capability and information content as knowledge evolves. The Tensor Net function reaches extremal values when the difference between decision probability and entropy ($\mathbf{D}_k^{(l)} - \mathbf{H}_k^{(l)}$) is maximized or minimized relative to changes in knowledge. Points where $\mathbf{D}_k^{(l)} = \mathbf{H}_k^{(l)}$ are significant as they represent transitions in the information processing dynamics. At these points, the instantaneous contribution to the Tensor Net function becomes zero, marking a critical boundary in the information flow.

Figure 8 illustrates the evolution of the Tensor Net function across all four layers during training. Several remarkable patterns emerge:

- **Zero-to-Zero Evolution:** All layers' Tensor Net functions begin at zero and eventually return to zero, despite following dramatically different trajectories.
- **Layer-Specific Dynamics:** Layer 2 (orange) exhibits a strong positive excursion reaching approximately 2.5×10^7 around step 16, while Layer 3 (green) shows a negative excursion bottoming at about -0.5×10^7 near step 23. Layer 1 (blue) displays a modest positive peak, and Layer 4 (red) remains close to zero throughout training.
- **Synchronized Convergence:** Despite their different magnitudes and opposite signs, all layers' Tensor Net functions converge to zero at approximately the same time (around step 45), suggesting a coordinated equilibrium state across the network.




Figure 7: Evolution of the Tensor Net function across layers during single-pass SKA training. Despite different trajectories and magnitudes, all layers approach zero around step 45-50, indicating completion of structured knowledge accumulation.

The zeroing of the Tensor Net function provides a principled criterion for determining when structured knowledge accumulation is complete. When the Tensor Net function approaches zero for a layer, it indicates that the layer has reached an information-theoretic equilibrium where:

$$\int \text{entropy } d\mathbf{Z} = \int \text{decision_probability } d\mathbf{Z} \quad (31)$$

This equilibrium represents the optimal balance between entropy and decision capability—the point at which the layer has accumulated sufficient knowledge to support its decision-making function without encoding redundant or irrelevant information. Unlike traditional stopping criteria based on arbitrary thresholds or epoch counts, this equilibrium emerges naturally from the system dynamics.

Importantly, this zeroing occurs even as the absolute entropy values continue to evolve (as shown in Figure 2), indicating that the Tensor Net function captures a more fundamental property of information dynamics than raw entropy measurements. The synchronized convergence across layers suggests a global coordination in the network despite each layer operating with local information only.

The Tensor Net function thus serves as both a theoretical insight into the nature of structured knowledge and a practical tool for monitoring learning progress in SKA networks. When all layers reach this equilibrium state, the network has completed its knowledge structuring phase and achieved a balanced distribution of information across its hierarchical representation.

9.7 Theoretical Interpretation

The observed entropy equilibrium aligns with the SKA principle that **knowledge alignment drives learning**. The minimization of entropy at each layer leads to a balance where knowledge accumulation is harmonized across layers, resulting in a structured learning process.

This behavior suggests the existence of a fundamental law in SKA-based neural networks:

In an SKA neural network, layer-wise entropy evolves towards an equilibrium state, where knowledge accumulation stabilizes across hierarchical representations. This equilibrium state is characterized precisely by the Tensor Net function criterion, where the instantaneous Tensor Net function reaches zero, marking the completion of structured knowledge accumulation and optimal balance between entropy and decision-making capability.

10 Implications and Future Work

The SKA framework offers transformative potential for neural network design and application, rooted in its continuous entropy reduction principles. Key implications and research directions include:

- **Autonomous Learning:** Layers independently minimize entropy, enabling self-sufficient training without global coordination.
- **Decentralization:** Local updates allow distribution across separate hardware, enhancing scalability in parallel systems.
- **Efficiency:** Forward-only learning reduces memory demands compared to backpropagation-based methods.
- **Interpretability:** Monitoring $\cos(\theta_k^{(l)})$ provides insight into layer-wise knowledge alignment.
- **Experiments:** Future work will compare SKA’s performance to stochastic gradient descent (SGD) on datasets like MNIST, validating its efficacy.
- **Real-Time Processing:** Single-pass updates suit SKA for processing live data streams efficiently.

11 Conclusion

The Structured Knowledge Accumulation (SKA) framework represents a fundamental paradigm shift in how we understand and implement neural learning systems. By reimagining entropy as a continuous process of knowledge accumulation expressed as:

$$H = -\frac{1}{\ln 2} \int z dD, \quad (32)$$

we have established a mathematical foundation that derives the sigmoid function as an emergent property and enables learning without backpropagation.

The framework’s core strengths emerge from several key innovations:

1. **Layer-wise entropy reduction** through local optimization, creating a biologically plausible learning mechanism that eliminates the need for error backpropagation.
2. **Self-organizing knowledge structures** that consistently follow similar entropy trajectories regardless of initialization, converging toward equilibrium states across layers.
3. **Progressive knowledge compression** in deeper layers, with entropy minima shifting toward lower knowledge magnitudes, representing increasingly efficient representations.
4. **Scalable, decentralized architecture** that enables distributed processing with significantly reduced computational requirements compared to traditional approaches.
5. **Continuous adaptation** to streaming inputs, making it ideal for real-time applications in computer vision, audio processing, and other domains.

The mathematical elegance of the tensor-based formulation, combined with the framework’s practical efficiency, positions SKA as both theoretically significant and practically applicable. By demonstrating effective learning on classification tasks with limited samples and forward-only processing, we have shown that entropy minimization itself may be a sufficient organizing principle for intelligent systems.

Beyond its immediate applications in machine learning, SKA offers a new lens through which to understand how information naturally structures itself in hierarchical systems. This could have profound implications for fields ranging from neuroscience to complex systems theory, potentially providing insights into how biological intelligence emerges from local information processing.

As we move forward with real-time experiments in visual and audio domains, we anticipate discovering new applications for this “plug and learn” approach that can adapt continuously to environmental inputs. The framework’s ability to discover hidden patterns through its fundamental entropy reduction principle may accelerate knowledge discovery across multiple scientific disciplines.

SKA represents not merely an incremental improvement to existing neural network methodologies, but a reconceptualization of how knowledge accumulates and organizes within intelligent systems. By bridging information theory and neural learning through a continuous, dynamic perspective on entropy, we open new avenues for developing more efficient, transparent, and biologically plausible artificial intelligence.

References

- [1] C. E. Shannon, “A Mathematical Theory of Communication,” *Bell System Technical Journal*, vol. 27, pp. 379–423, 1948.